

The *Compliant* Plugin

François Faure

2012

1 Introduction

This plugin serves two purposes:

- to provide a new, more complete mechanism for assembling the dynamical system matrices such as mass \mathbf{M} and stiffness \mathbf{K} , as well as constraint Jacobians \mathbf{J} ,
- and experiment a unified approach of soft and hard constraints, inspired from [?].

Contrary with other implementations of matrix assembly available in SOFA, this one handles mappings, and even multimappings.

Sections ?? and ?? explain the unified constraint approach, while Section ?? presents the matrix assembly process based on an example. An overview of the corresponding API is given in Section ??.

2 Constraint forces

This approach unifies soft and hard constraints, by providing constraints with compliance.

2.1 Hard constraints

For a good introduction on hard constraints, see [?]. Hard constraints are usually implemented using Lagrange multipliers λ in the following equation:

$$\begin{pmatrix} \mathbf{M} & -\mathbf{J}^T \\ \mathbf{J} & \end{pmatrix} \begin{pmatrix} \mathbf{a} \\ \lambda \end{pmatrix} = \begin{pmatrix} \mathbf{f}_{ext} \\ -\ddot{\phi} \end{pmatrix} \quad (1)$$

where \mathbf{M} is the mass matrix (or, more generally, a dynamics matrix such as $\mathbf{M} - h^2\mathbf{K}$ used in implicit time integration), \mathbf{J} is the Jacobian matrix of the constraint(s), \mathbf{a} is the acceleration, \mathbf{f}_{ext} is the net external force applied to the system, λ is the constraint force, ϕ is the constraint violation and $\ddot{\phi}$ is the second time derivative of the violation (i.e. the error on accelerations). This form of equation system is typically called *Karush-Kuhn-Tucker (KKT)*. λ and $\ddot{\phi}$ are vectors with as many entries as scalar constraints. The equation system is typically solved using a Schur complement to compute the constraint forces:

$$\mathbf{J}\mathbf{M}^{-1}\mathbf{J}^T\lambda = -\ddot{\phi} - \mathbf{J}\mathbf{M}^{-1}\mathbf{f}_{ext} \quad (2)$$

and then the acceleration is computed as $\mathbf{a} = \mathbf{M}^{-1}(\mathbf{f}_{ext} + \mathbf{J}^T\lambda)$.

2.2 Projective constraints

When the constraints are simple enough, such as keeping points fixed or in some hyperplane, they are more easily handled by straightforwardly projecting forces and displacements. The previous equation becomes:

$$\begin{pmatrix} \mathbf{P}\mathbf{M} & -\mathbf{P}\mathbf{J}^T \\ \mathbf{J}\mathbf{P} & \end{pmatrix} \begin{pmatrix} \mathbf{a} \\ \lambda \end{pmatrix} = \begin{pmatrix} \mathbf{P}\mathbf{f}_{ext} \\ -\dot{\phi} \end{pmatrix} \quad (3)$$

where \mathbf{P} is the orthogonal projection matrix corresponding to the constraint.

2.3 Generalized constraints

In the generalized approach, the constraint forces are considered directly proportional to the constraint violations ϕ and their first time derivative (error on velocity) $\dot{\phi}$:

$$\lambda = -\frac{1}{c} (\phi + d\dot{\phi}) \quad (4)$$

where the positive real number c is the compliance of the constraint and d is its damping ratio. Combined with a time discretization scheme, this leads to an equation system similar to (??) as shown in Section ??.

2.4 Force or constraint ?

ForceFields usually contribute the top line of Equation ??, by accumulating force in the right-hand term. In implicit integration, their stiffness matrix is also accumulated to the left-hand side. When a ForceField has an invertible stiffness matrix, it can alternatively be handled as a constraint with non-null compliance, rather than a standard force. In this case, it contributes to the bottom of Equation ?? rather than to the top.

Due to matrix conditioning issues, and depending on which linear solver is used, it may be a good idea to process very stiff forces as low compliances rather than large stiffnesses. In this case, we call them *compliant* or *compliance forces*. Here the term compliant does not convey information on the magnitude of the stiffness (or compliance) but only on the way it is processed by the solver, namely in the bottom row of the KKT equation system. Note, however, that each force handled as a constraint increases the size of the equation system, since it adds lines to the bottom (and columns to the right) of the equation system. Conversely, it may be more efficient to handle soft forces as stiffnesses, since the size of the stiffness matrix depends on the number of independent DOFs, which does not increase along with the number of forces. We call such forces *stiff* or *stiffness forces*, though here again, the term does not convey information on the magnitude of the stiffness but only on the way it is handled by the solver, namely in the first row of the KKT equation system.

Note that geometric stiffness is not (yet) handled in the compliance view, which may result in instabilities in case of large displacements.

3 Time integration

Our integration scheme is:

$$\mathbf{v}_{n+1} = \mathbf{v}_n + h\mathbf{M}^{-1}(\alpha\mathbf{f}_{n+1} + (1-\alpha)\mathbf{f}_n) \quad (5)$$

$$\mathbf{x}_{n+1} = \mathbf{x}_n + h(\beta\mathbf{v}_{n+1} + (1-\beta)\mathbf{v}_n) \quad (6)$$

where index n denotes current values while index $n+1$ denotes next values, α is the implicit velocity factor, and β is the implicit position factor. Let

$$\Delta\mathbf{v} = \mathbf{v}_{n+1} - \mathbf{v}_n = h\mathbf{M}^{-1}(\alpha\mathbf{f}_{n+1} + (1-\alpha)\mathbf{f}_n) \quad (7)$$

$$\Delta\mathbf{x} = \mathbf{x}_{n+1} - \mathbf{x}_n = h(\beta\Delta\mathbf{v} + \mathbf{v}_n) \quad (8)$$

be the velocity and position changes across the time step. We have not carefully studied the influence of these parameters, but it seems that $\alpha = 0.5$ and $\beta = 1$ corresponds to an **energy-conserving integration scheme**.

The constraint violation ϕ and its Jacobian \mathbf{J} are:

$$\mathbf{J} = \frac{\partial\phi}{\partial\mathbf{x}} \quad (9)$$

$$\phi_{n+1} \simeq \phi_n + \mathbf{J}\Delta\mathbf{x} = \phi_n + h\dot{\phi} + \mathbf{J}h\beta\Delta\mathbf{v} \quad (10)$$

$$\dot{\phi}_{n+1} \simeq \dot{\phi}_n + \mathbf{J}\Delta\mathbf{v} \quad (11)$$

The corresponding forces are:

$$\mathbf{f} = \mathbf{f}_{ext} + \mathbf{J}^T\lambda \quad (12)$$

$$\lambda_i = -\frac{1}{c_i}(\phi_i + d\dot{\phi}_i) \quad (13)$$

where the subscript i denotes a scalar constraint.

The average constraint forces are computed using equations ??, ?? and ??:

$$\begin{aligned} \bar{\lambda}_i &= \alpha\lambda_{n+1} + (1-\alpha)\lambda_n \\ &= -\frac{1}{c_i}(\alpha\phi + \alpha h\dot{\phi} + \alpha h\beta\mathbf{J}\Delta\mathbf{v} + \alpha d\dot{\phi} + \alpha d\mathbf{J}\Delta\mathbf{v} + (1-\alpha)\phi + (1-\alpha)d\dot{\phi}) \\ &= -\frac{1}{c_i}(\phi + d\dot{\phi} + \alpha h\dot{\phi} + \alpha(h\beta + d)\mathbf{J}\Delta\mathbf{v}) \end{aligned}$$

We can rewrite the previous equation as:

$$\mathbf{J}\Delta\mathbf{v} + \frac{1}{\alpha(h\beta + d)}\mathbf{C}\bar{\lambda} = -\frac{1}{\alpha(h\beta + d)}(\phi + (d + \alpha h)\dot{\phi}) \quad (14)$$

where values without indices denote current values. The complete equation system has the following generalized KKT form:

$$\begin{pmatrix} \frac{1}{h}\mathbf{P}\mathbf{M} & -\mathbf{P}\mathbf{J}^T \\ \mathbf{J} & \frac{1}{l}\mathbf{C} \end{pmatrix} \begin{pmatrix} \Delta\mathbf{v} \\ \bar{\lambda} \end{pmatrix} = \begin{pmatrix} \mathbf{P}\mathbf{f}_{ext} \\ -\frac{1}{l}(\phi + (d + \alpha h)\dot{\phi}) \end{pmatrix} \quad (15)$$

where $l = \alpha(h\beta + d)$. The matrix in equation ?? is semi-definite when matrix \mathbf{C} is null, or when matrix \mathbf{P} removes independent degrees of freedom without reducing the size of the mass matrix accordingly. This equation system can be




Figure 1: Example of mechanical system structure. The X represent mechanical states, while the arrows represent the kinematic hierarchy, and the plain lines represent components acting on the states.

solved using iterative solvers, such as CG or other Krylov methods. An other option is to create a reduced equation system on constraint forces only:

$$\left(h\mathbf{J}\bar{\mathbf{M}}^{-1}\mathbf{J}^T + \frac{1}{l}\mathbf{C} \right) \bar{\lambda} = -\frac{1}{l} \left(\phi + (d + h\alpha)\dot{\phi} \right) - h\mathbf{J}\mathbf{P}^T\bar{\mathbf{M}}^{-1}\mathbf{f}_{ext} \quad (16)$$

where matrix $h\mathbf{J}\bar{\mathbf{M}}^{-1}\mathbf{J}^T + \frac{1}{l}\mathbf{C}$ is called the Schur complement of the KKT equation system. Matrix $\bar{\mathbf{M}} = \mathbf{P}\mathbf{M}\mathbf{P}^T + \text{diag}(\epsilon)$ is optionally regularized by the linear solver using a very small number ϵ . This makes the matrix PSD, and does not significantly change the value of the equation solution thanks to the small magnitude of the perturbation.

Solving for $\bar{\lambda}$ allows us to update the velocities and the positions:

$$\Delta \mathbf{v} = h\mathbf{P}\bar{\mathbf{M}}^{-1}(\mathbf{f}_{ext} + \mathbf{J}^T\bar{\lambda}) \quad (17)$$

$$\Delta \mathbf{x} = h(\mathbf{v} + \beta\Delta \mathbf{v}) \quad (18)$$

4 Matrix assembly

The equation system, in its most general form, can be written as:

$$\begin{pmatrix} \mathbf{M} & -\mathbf{J}^T \\ \mathbf{J} & \mathbf{C} \end{pmatrix} \begin{pmatrix} \mathbf{x} \\ \lambda \end{pmatrix} = \begin{pmatrix} \mathbf{f} \\ \phi \end{pmatrix} \quad (19)$$

We assemble the 7 terms of equation ?? separately.

Figure ?? shows an example of mechanical system. The independent DOFs are X_a and X_d . State X_b is attached to X_a using a simple mapping, and a mass matrix M_{bb} is defined at this level. State X_c is attached to X_b using a simple mapping, and a compliance matrix $C_{\alpha\alpha}$ (possibly a deformation force) is applied to these DOFs. State X_e is attached to X_a and X_d at the same time, using a MultiMapping. A compliance matrix $C_{\beta\beta}$, possibly an interaction force, is applied to these DOFs, while a mass M_{dd} is applied to X_d .

The corresponding equation system has the block structure shown in the right of Figure ?. The main blocks of the equation system are highlighted in grey rectangles. The J matrices are the mapping matrices. The bottom row has two mappings, since the state X_e impacted by compliance β depends on two parent states.

The figure shows a block matrix equation. On the left, a large square matrix is partitioned into four quadrants. The top-left quadrant contains a yellow block labeled $J_{ba}^T M_{ab} J_{ba}$. The top-right quadrant is white and labeled $(-sym)$. The bottom-left quadrant is white and contains two yellow blocks: $J_{cb} J_{ba}$ in the top-left and J_{ea} in the bottom-left, with a yellow block J_{ed} in the bottom-right. The bottom-right quadrant contains two yellow blocks: $C_{\alpha\alpha}$ in the top-left and $C_{\beta\beta}$ in the bottom-right, with a white block M_{dd} in the top-right. To the right of the matrix is an equals sign, followed by two vertical vectors. The first vector has four yellow blocks labeled x_a , x_d , λ_{α} , and λ_{β} . The second vector has four yellow blocks labeled f_a , f_d , φ_{α} , and φ_{β} .

Figure 2: Block view of the equation ?? applied to the system of Figure ??, with non-null blocks highlighted in yellow.

The assembly of local matrices and vectors to the global matrices and vectors which compose the system is performed using shift matrices, as illustrated in Figure ???. Shift matrices J_{*0} , shown in the top of the figure, are composed of identity matrices (represented with a diagonal in a block) and null blocks. They can be used to implement the shifting of vector entry indices from local to global. The assembly is thus easily expressed as a sum of product of matrices, as illustrated in Figure ???. While shifting values in dense arrays using matrix products is not efficient (vector assembly is actually implemented using shifted copies), the sum of matrix products is a reasonable implementation of sparse matrix assembly.

5 Implementation

5.1 Dependencies

This plugin uses **Eigen** and **suitesparse**. Moreover, the test suite requires the **boost unit test framework**, see <https://wiki.sofa-framework.org/wiki/UnitTesting>.

5.2 Functions

The virtual functions used for setting up the equation system are declared in `BaseForceField.h`, `BaseMapping.h`, `BaseMass.h`, `BaseProjectiveConstraint.h`, `MechanicalParams.h`. They are documented in “Experimental Compliance API” documentation blocks of these files.

5.3 BaseForceField functions

ForceFields may be handled as forces or as soft constraints, as explained in Section ???. The default behavior is to handle them as forces, and in this case their force is accumulated using a classical `MechanicalOperations::computeForce`




Figure 3: Shift matrices J_{*0} are used to place blocks within global vectors and matrices.

$$\begin{array}{cc|cc}
 \begin{array}{c} J_{a0}^T J_{ba}^T M_{bb} J_{ba} J_{a0} \\ + \\ J_{d0}^T M_{dd} J_{d0} \end{array} & & & & \begin{array}{c} x_a \\ x_d \end{array} & \begin{array}{c} f_a \\ f_d \end{array} \\
 & & (-\text{sym}) & & & \\
 \begin{array}{c} J_{\alpha 0}^T J_{cb} J_{ba} J_{a0} \\ + \\ J_{\beta 0}^T J_{ea} J_{a0} + J_{\beta 0}^T J_{ed} J_{d0} \end{array} & & \begin{array}{c} J_{\alpha 0}^T C_{\alpha\alpha} J_{\alpha 0} \\ + \\ J_{\beta 0}^T C_{\beta\beta} J_{\beta 0} \end{array} & & \begin{array}{c} \lambda_\alpha \\ \lambda_\beta \end{array} & \begin{array}{c} q_\alpha \\ q_\beta \end{array} \\
 & & & & = &
 \end{array}$$

Figure 4: The assembly is performed by summing the global vectors and matrices obtained by shifting local vectors and matrices.

operation, while their stiffness matrix is accumulated using the `getStiffnessMatrix` briefly described below.

- `getComplianceMatrix` returns NULL if the ForceField is to be handled as a traditional force function, while it returns a pointer to a compliance matrix if it is to be handled as a constraint.
- `addForce` is used to accumulate the force in the right-hand term. This function is not specific to the *Compliant* plugin. This should do nothing when the force must be handled as a constraint.
- `getStiffnessMatrix` is the complement of `getCompliantMatrix`. If one returns NULL, then the other must return a pointer to a matrix.
- `writeConstraintValue` is used to write the constraint violation ϕ in Eq.?? when the ForceField is handled as a constraint.
- `getDampingRatio` is used in the constraint case, and corresponds to parameter d in Eq.??

5.3.1 BaseMapping functions

- `getJs` returns a list of Jacobian matrices, one for each parent of the mapping. Typical mappings have only one parent, but MultiMappings have several ones.
- `getKs` returns a list of stiffness matrices, one for each parent. These correspond to geometric stiffness, i.e. change of forces due to mapping nonlinearity, like `addDJT` in the traditional API. **TODO:** refine this, to also return off-diagonal terms in MultiMappings.

5.3.2 BaseMass functions

- `addMToMatrix` to create the mass matrix. This function is not specific to the *Compliant* plugin.

5.4 BaseProjectiveConstraint functions

- `projectMatrix` replaces a matrix with one corresponding to the same linear operator, projected to cancel, out the forbidden directions.

5.5 MechanicalParams functions

- `implicitVelocity` is the α parameter in the implicit velocity integration of Eq.??.
- `implicitPosition` is the β parameter in the implicit position integration of Eq.??.

6 Alternative model

To allow more flexibility in the definition of damping in the compliance approach, the following variant is possible. We replace eq. ?? with the following.

$$\lambda = -\mathbf{C}^{-1}\phi - \mathbf{D}\dot{\phi} \quad (20)$$

where ϕ and $\dot{\phi}$ are vectors, and \mathbf{C} and \mathbf{D} symmetric matrices. This allows us to define non-uniform damping coefficients. We rewrite eq. ?? as:

$$\alpha(h\beta\mathbf{I} + \mathbf{CD})\mathbf{J}\Delta\mathbf{v} + \mathbf{C}\bar{\lambda} = -\phi - \alpha(h\mathbf{I} + \mathbf{CD})\dot{\phi} \quad (21)$$

where \mathbf{I} is the identity matrix of the appropriate size. Unfortunately, the diagonal matrix in the left of \mathbf{J} in the term on $\Delta\mathbf{v}$ now makes the Schur complement unsymmetric. Therefore, we have not implemented this approach.

7 Additional Notes

The initial implementation of the compliance framework is provided by `ComplianceSolver`, which performs system assembly, system solves and time-stepping. It uses a sparse Cholesky solver and is able to handle bilateral constraints only.

With time, this component has been broken down to the following components:

- `CompliantImplicitSolver`, which performs system assembly (through an internal `AssemblyVisitor`) and time-stepping.
- `KKTSolver`, which performs the KKT solve associated with an `AssembledSystem`, itself provided by the `CompliantImplicitSolver` through its internal `AssemblyVisitor`.

`KKTSolver` is a virtual base class implemented in `MinresSolver` and `LDLTSolver` for bilateral constraints. The `CompliantDev` plugin contains additional KKT solvers for unilateral and frictional constraints.

We now give a quick user documentation for the above components.

7.1 CompliantImplicitSolver

This component implements a linearized implicit Euler solver. The KKT system is first assembled by an `AssemblyVisitor`, and factorized by the `KKTSolver`.

The right-hand-side is then computed depending on constraint stabilization and order of resolution, and the `KKTSolver` provides system solves.

At the end of the time-step, velocities are updated and the system is stepped forward. Optionally, one can propagate the computed Lagrange multipliers upwards in the scene graph at the end of the time step in the *force* state vector for easy access.

7.1.1 Order of Resolution

`use_velocity="true, false"`

Formulates the system at the velocity level. Acceleration is used otherwise. This is mostly for debugging purpose, unless warm-starts are disabled (see below).

7.1.2 Warm Starts

`warm_start="true, false"`

Warm starts iterative solvers with previous velocity/lambda (velocity-level). At the acceleration level, only the lambda are used.

When disabled and used with iterative solvers and premature exit, this biases the solution towards zero, thus this results in numerical damping with velocity-level.

This is mostly for debugging purpose.

7.1.3 Constraint Stabilization

`stablization="false, true"`

Performs simple constraint stabilization. This causes an additional KKT solve during the time step.

You need to add a **Stabilization** component next to any compliance you wish to stabilize. Only use with zero-compliance (holonomic constraints).

TODO more, example

7.2 Lagrange Multiplier Propagation

`propagate_lambdas="false, true"`

Optionally propagate/aggregate Lagrange multipliers from compliant DOFs to independent DOFs at the end of the time step. This erases the *force* state vector.

TODO example

7.3 LDLTSolver

no options

TODO more

7.4 MinresSolver

`iterations="false, true"`

iteration count

`precision="1e-8"`

residual norm threshold

`relative="true, false"`

use relative precision, otherwise absolute

TODO parallel ? TODO schur complement ?

TODO better description